

Frequently Bought Together

...and the three things it can't do

60 minutes · illustrative portfolio scenario · simulated telemetry · nothing deployed

AI Engineering · business engagement

Would you have shipped this?

A row from the simulated adoption table - every number below is generated test data

WHITE METAL LANTERN → WHITE HANGING HEART T-LIGHT HOLDER

63 baskets, confidence 0.28, lift 2.70 — looks like a finding

...until you notice **the same t-light holder is recommended for 69 other products**

It is in a large share of every basket, so it pairs with everything.

Hold that thought

Every number here is from the live table.

By 0:30 you'll be able to say in one sentence why this row is weak.

Popularity is not affinity. Always ask what else that product is recommended for.

Reading a recommendation

One real row, one number at a time

CHILDS GARDEN SPADE PINK → CHILDS GARDEN SPADE BLUE

40

baskets

Your sample size. Read this first — a huge lift on 8 baskets is noise.

0.85

confidence

Of baskets with the pink spade, 85% also had the blue one.

234

lift

234× more likely than chance. This is the ranking column.

Back to the first row: lift 2.7 against 234 here. Same catalogue, two very different kinds of evidence.

Two signals — and they are not equal

Check the method column before you build a plan on a row



co_purchase

Built from what people actually bought together.

The real signal — trust it in proportion to the basket count.

Covers about 30% of the catalogue.



content_tfidf

Built from the product description text.

Nobody has been observed buying these together — they simply read alike.

A starting point, never evidence.

Why keep the weaker signal? Without it 7 products in 10 show an empty panel — and people generalise from the first empty screen they see.

Three things it cannot do

Said first and unprompted — the credibility of everything else depends on it

1 It cannot tell you why

Bunting and paper plates co-occur because both are party purchases — not because one drives the other. Placing them together is sensible. Discounting one to lift the other is not.

2 It cannot see what you didn't stock


A product out of stock for six weeks builds its pairings from the weeks it was available. It will look less connected than it is.

3 It cannot decide anything

Nothing writes to a trading system. Every recommendation is a suggestion a person accepts or rejects. That is a design decision, not a limitation.

How it's going, and one ask

Simulated weekly reach — % of 62 fictional users active that week



Weeks 3 and 8 are these workshops

16% → 34% after the first.

27% → 44% after the second.

Both held rather than decaying — which is the part that matters. A spike that decays means people looked once.

One ask: use it once this week on a real decision, and tell us what happened. The feedback button is two clicks.